

Parallel Numerics

Exercise 3: Sparse Matrices

1 Sparse Matrices and Graphs

Given the matrix

$$A = \begin{pmatrix} 4 & 0 & 7 & 0 & 0 & 0 \\ 0 & 5 & 42 & 0 & 12 & 3 \\ 0 & 4 & 7 & 2 & 0 & 5 \\ 0 & 0 & 44 & 32 & 0 & 0 \\ 0 & 0 & 0 & 0 & 99 & 8 \\ 100 & 2 & 0 & 0 & 0 & 102 \end{pmatrix}$$

- i) What methods we know to store sparse matrices? What are advantages and disadvantages? Show how to store matrix A using the various formats.
- ii) What graph datastructure can be used to represent the sparsity pattern of a (non-) symmetric matrix? How many nodes are needed?
- iii) What is an adjacency matrix? How can it be obtained? How does it relate to the beforementioned graph?
- iv) Give the graph of matrix A . Based on the properties of the matrix, what kind of graph you need?
- v) How can not only the adjacency of a matrix be saved in a graph, but also the values?

2 Parallel QR

We consider tall and skinny QR as used, e.g. in the first step of block-wise parallel QR decomposition.

$$\overbrace{\begin{bmatrix} A \end{bmatrix}}^m = \begin{bmatrix} Q \end{bmatrix} \cdot [R], \quad A \in \mathbb{R}^{n \times m}$$

based on the Gram-Schmidt orthonormalization as described in the lecture notes in 4.2. For parallelization, we use a decomposition of A und Q into p row-blocks. Processor i in a distributed memory system stores A_i and Q_i . R is stored at the master process ($i = 1$).

$$\begin{bmatrix} A_1 \\ A_2 \\ \vdots \\ A_p \end{bmatrix} = \begin{bmatrix} Q_1 \\ Q_2 \\ \vdots \\ Q_p \end{bmatrix} \cdot [R]$$

i) Calculate the communication costs, i.e.

- a) The number of messages and
- b) the amount of data (number of matrix entries)

for each of the five algorithmic steps and for the whole QR-decomposition of A .

ii) Compare i with the communication avoiding QR-algorithm for $p = 2^q$, which calculates small QR-decomposition for the blocks of A independently and combines them to larger blocks in a fan-in-like process

$$\begin{aligned} & \begin{bmatrix} A_1 \\ A_2 \\ A_3 \\ A_4 \end{bmatrix} \Rightarrow \\ & \begin{array}{l} p=1: A_1 = Q_1 R \\ p=2: A_2 = Q_2 R \\ p=3: A_3 = Q_3 R \\ p=4: A_4 = Q_4 R \end{array} \Rightarrow \begin{array}{l} p=1: \begin{pmatrix} Q_1 & & & \end{pmatrix} \begin{pmatrix} R_1 \\ R_2 \\ R_3 \\ R_4 \end{pmatrix} = \begin{pmatrix} Q_1 & & & \end{pmatrix} \begin{bmatrix} \bar{Q}_1 \\ \bar{Q}_3 \end{bmatrix} \begin{bmatrix} \bar{R}_1 \\ \bar{R}_3 \end{bmatrix} \\ p=3: \begin{pmatrix} & Q_2 & & \\ Q_3 & & & \\ & & Q_4 & \end{pmatrix} \begin{pmatrix} R_1 \\ R_2 \\ R_3 \\ R_4 \end{pmatrix} = \begin{pmatrix} & Q_2 & & \\ Q_3 & & & \\ & & Q_4 & \end{pmatrix} \begin{bmatrix} \bar{Q}_1 \\ \bar{Q}_3 \end{bmatrix} \begin{bmatrix} \bar{R}_1 \\ \bar{R}_3 \end{bmatrix} \end{array} \Rightarrow \\ & \begin{array}{l} p=1: \begin{pmatrix} Q_1 & & & \\ & Q_2 & & \\ & & Q_3 & \\ & & & Q_4 \end{pmatrix} \begin{pmatrix} \bar{Q}_1 \\ & & & \\ & \bar{Q}_3 & & \end{pmatrix} \begin{pmatrix} \bar{R}_1 \\ \bar{R}_3 \end{pmatrix} = \\ \underbrace{\begin{pmatrix} Q_1 & & & \\ & Q_2 & & \\ & & Q_3 & \\ & & & Q_4 \end{pmatrix} \begin{pmatrix} \bar{Q}_1 \\ & & & \\ & \bar{Q}_3 & & \end{pmatrix} \begin{bmatrix} \bar{Q}_1 \\ \bar{Q}_3 \end{bmatrix}}_Q \begin{bmatrix} R \end{bmatrix} \end{array} \end{aligned}$$

iii) Implement both variants, compare scalability and runtimes for $m = 100, 3200$ and $p = 1, 2, 4, 8, 16, 32$.